



# Design and Implementation of a Deterministic Multipath Solution for Optimizing Backhaul Performance in Geographically Distributed 5G Campus Networks

Nicolas Zunker

n.zunker@campus.tu-berlin.de

December 7, 2023

MASTER'S THESIS

Telecommunication Systems Institute

Technische Universität Berlin

Examiner 1: Prof. Dr. Thomas Magedanz

Advisor: Dr. Ing. Marius Corici

Examiner 2: Prof. Dr. Axel Küpper



# Declaration

I hereby declare that I have created the present work independently and by my own without illicit assistance and only utilizing the listed sources and tools.

Hiermit erkläre ich, dass ich die vorliegende Arbeit selbstständig und eigenhändig sowie ohne unerlaubte fremde Hilfe und ausschliesslich unter Verwendung der aufgeführten Quellen und Hilfsmittel angefertigt habe.

Die selbständige und eigenständige Anfertigung versichert an Eides statt:

Berlin, December 7, 2023

Nicolas Zunker



# Acknowledgment

I would like to thank my unofficial advisor Hauke Buhr for his help, advice and support. I also thank my advisor Marius for his insights. Both of them provided excellent constructive feedback to me, for which I am grateful. I am also indebted to the rest team at Fraunhofer FOKUS for all their help, especially Hemant Zope, Eric Troudt, Christian Scheich and Aleksandar Yonchev. Lastly I would like to thank Prof. Dr.-Ing. habil. Falko Dressler and Prof. Dr. Zubow whose courses I enjoyed immensely during my studies, and from whom I learned a lot.

NZ



# Abstract

TODO

# Table of Contents

Declaration . . . . .	I
Acknowledgment . . . . .	III
Abstract . . . . .	V
List of Figures . . . . .	IX
<b>1 Introduction</b>	<b>1</b>
1.1 Motivation and Problem Description . . . . .	1
1.2 Goal . . . . .	2
1.3 Structure . . . . .	3
<b>2 Background and Related Work</b>	<b>5</b>
2.1 Mobile Networks and Backhaul . . . . .	5
2.1.1 Backhaul in 5G . . . . .	5
2.1.2 5G and Campus Networks . . . . .	5
2.2 Multipathing and Multihoming . . . . .	5



2.2.1	Collecting Per-Link Metrics . . . . .	5
2.2.2	How to Guarantee QoS . . . . .	7
2.3	Determinism in Computer Networking . . . . .	8
2.3.1	Quality of Service (QoS) . . . . .	8
2.3.2	Deterministic Networking Specification . . . . .	8
2.3.3	Traffic Shaping . . . . .	8
<b>3</b>	<b>Approach</b>	<b>9</b>
3.1	Skeleton Structure . . . . .	9
3.2	Architectural Components Required for Deterministic Multipath Backhaul in a 5G Campus Environment . . . . .	10
3.2.1	Elimination of Contention Loss . . . . .	11
3.2.2	Jitter Reduction and Latency Guarantees . . . . .	12
3.2.3	Service Protection . . . . .	13
3.2.4	Multipath Considerations . . . . .	16
3.2.5	Summary of Requirements . . . . .	17
3.3	Overview of the WAN Connector's Features and Components . . . . .	17
3.3.1	Path Selection Algorithm . . . . .	17
3.3.2	Packet Ordering and De-Duplication Function . . . . .	19
3.3.3	Internal Architecture . . . . .	20

3.3.4 Overview . . . . .	23
<b>4 Evaluation</b>	<b>25</b>
4.1 Approach to Evaluating Performance . . . . .	25
4.2 Latency Based Path Switching . . . . .	27
4.3 Reliability Based Path Switching . . . . .	27
4.4 Extended Experiments . . . . .	27
<b>5 Conclusion and Outlook</b>	<b>29</b>
5.1 Improvements . . . . .	29
5.2 Implications and Further Areas of Research . . . . .	30
<b>Bibliography</b>	<b>33</b>

# List of Figures

1.1	5G Deployment with 2 UPFs . . . . .	2
3.1	Internal Architecture of the WAN Connector . . . . .	23
4.1	Testbed Setup . . . . .	26



# Chapter 1

## Introduction

### 1.1 Motivation and Problem Description

One of the aims for the fifth generation of mobile networks (5G) and its successors will be a greater diversification of the classes of service. As the use cases for these networks evolve, there is a greater need for quality of service (QoS) tailored to each use case. For example, in the Industrial Internet of Things (IIoT) the requirements on latency, jitter, and reliability may be extremely stringent. Supporting these kinds of classes of service can be a challenge for mobile network operators (MNOs) and will require novel approaches to familiar problems, such as backhaul.

As there are more heterogeneous edge deployments and more campus networks, backhaul becomes more challenging, since many sites may not have access to optical fibre, and may be forced into using other solutions such as satellite links, mmWave backhaul, or pre-existing on-site ISP connections. Providing the kind of deterministic quality of service that these sites may require can be a very difficult challenge.

Particularly with the rise of satellite backhaul options, fuelled by the new space race, many

remote deployments may choose to integrate satellite backhaul because fibre is infeasible or too far away in the future. These connections are usually being added in addition to existing one, and thus network operators may choose to utilize more than one backhaul connection at the same time. Either in order to increase the available bandwidth or to utilize the different qualities of the backhaul links. This bears the question whether multipathing could then be used to provide deterministic backhaul by intelligently selecting on which links to forward packets. This approach bears similarity to multihoming as well as to multi-path routing in Wireless Sensor Networks (WSNs), and can take inspiration from the existing body of research in these fields, which has demonstrated that QoS can be improved by using multiple links or paths simultaneously [1, 2, 17, 19, 22, 34].

## 1.2 Goal




Figure 1.1: 5G Deployment with 2 UPFs

The goal of this thesis is to design and provide an implementation of a Wide Area Network (WAN) connector, that can be placed at the ingress and egress point of two or more locations, and utilizes multipathing in order to provide deterministic backhaul between the two sites.

The performance of this approach will then be quantitatively analyzed in experiments.

Looking at Figure 1.1 we can see how this is envisioned to work. WAN connectors are deployed both in the geographically distributed 5G campus network, which has more than one egress link, and in the core. Then, using the multiple outgoing links, the traffic is backhauled to the other site, while respecting its QoS requirements. This can be especially beneficial for critical applications (e.g. between industrial sites), which are in locations that do not have access to optical fibre for backhaul.

For a 5G deployment the proposed WAN connector could be deployed in between two User Plane Functions (UPF) , in order to provide deterministic backhaul. The architecture for such a deployment is shown in Figure 1.1. Further, the on-site UPF is not a strict requirement; it is also feasible to connect the RAN directly to the WAN Connector.

## 1.3 Structure

This thesis will follow a simple 5 chapter structure. This section concludes the introduction chapter, what follows will be one chapter to provide both basic background information as well as to highlight the existing literature which is of relevance to the problem statement. Next will be the approach and evaluation chapters, where, respectively, the design of the implementation will be presented and discussed, and then its performance will be analyzed. The conclusion chapter will review the relevant findings, the successes and failures of the approach, and the possibilities for future research in this area.





# Chapter 2

## Background and Related Work

### 2.1 Mobile Networks and Backhaul

#### 2.1.1 Backhaul in 5G

#### 2.1.2 5G and Campus Networks

### 2.2 Multipathing and Multihoming

#### 2.2.1 Collecting Per-Link Metrics

##### Measurement-based Metrics

In [2] the authors collected both passive metrics (looking at response times for outgoing packets), and active measurements (sending ICMP ping, or TCP SYN messages and measuring the response time). Using the passive measurements enabled their multihomed approach to perform well, but when using the active measurements the performance was better. Crucially,

the passive measurements worked better over larger sampling periods, because it took longer to get a full overview of all the possible routes. Whereas the active sampling approach acquired its measurements faster and was thus more effective over smaller sampling intervals.

Considering these results, it is proposed to utilize both active and passive measurements. All three metrics- packet loss, latency, and jitter- will be periodically measured in an active manner. The period over which to perform these measurements is an important design decision for the WAN connector, and it will be met later.

Beyond this, these metrics will also be monitored on a passive basis wherever possible. In order to measure either the time needed (for latency and jitter), or to ascertain that a packet has been lost, a response is required for each outgoing message. This may only be possible for TCP's SYN and SYN ACK messages as well as other protocols which are guaranteed to contain request-response handshakes, and thus complicates the process of passive measurement.

For classical wired links in multihomed scenarios, [33] have observed that one link will generally dominate with regards to latency, but with brief periods where other links' performance is superior. These same authors also note that with regards to packet loss there is far less prevalence of a "dominant" link, and generally the links will perform comparably. Packet loss is also a particularly difficult metric to measure, since most links are highly reliable and when they do experience packet loss it is in bursts [33]. Wireless connections are usually less reliable and may experience more consistent rates of packet loss at the data link layer, however it is opaque from the perspective of the higher layers, which may only perceive it through the jitter and/or latency.

Ultimately, many of the best performing approaches for predicting packet loss, e.g. Hidden Markov Models [8, 33], are still somewhat imprecise and inaccurate. These models assume the link is in one of two states, good or bad, and each state has a different probability for packet loss, and there is a transition matrix which represents the probabilities of switching

from one state to another. This thesis will also use such a model to try to predict packet loss.

## **2.2.2 How to Guarantee QoS**

### **Limiting Jitter**

The design of an approach which guarantees QoS is also an interesting challenge. One idea to improve jitter when backhauling across multiple links is to duplicate packets and forward them on multiple links, and have the WAN connector on the other end buffer incoming packets and release them at a constant rate. This way, in the event of a packet being lost on one link, the other link is still able to receive it and the delay caused by retransmission is avoided. The downside of this approach is that it guarantees that the latency will always be as slow as the slowest link.

### **Low Delays**

For reducing latency it would appear likely that the simplest approach may be a greedy method (as in [17], in the online case) which always selects the lowest latency connection. However there is room for nuance here since the connection must not be overloaded and also because certain traffic may have very relaxed latency requirements but use up more bandwidth. This means monitoring the load on any one link will be important. Finally there are also more intelligent approaches, i.e. integer linear programming (used in [22], and used for the offline case in [17]) which find an optimal solution satisfying the given requirements.

The timescale over which to use a chosen link is also of interest. In [19] the time for which a link should be used is varied based on the predicted qualities of the link. These predictions are made based on past performance.

## **Error Rates**

Reliability presents yet another challenge. However in a multihomed scenario it becomes easier to guarantee this via duplication, and/or forward error correction (FEC). For example if a packet flow requires 99% reliability this can be guaranteed by duplicating packets across two links which are both only 90% reliable. Alternatively, in such a situation, an FEC configured for 10% packet loss could be used to pre-code the packets sent across one of the links, and thus increase the reliability to the required level.

Although duplication uses a lot of bandwidth, in order to support 5G's ultra-reliable low latency (URLLC) QoS requirements, which are especially relevant for IIoT applications, it may be the only option for certain traffic flows. Forward Error Correction is an excellent protection against consistently lossy links, however it may fail to be reliable when there are concentrated bursts of dropped packets, which is a more common occurrence in packet switched networks. Either way, the effectivity of FEC in a deterministic backhaul unit is an interesting question which this thesis may also explore.

## **2.3 Determinism in Computer Networking**

### **2.3.1 Quality of Service (QoS)**

### **2.3.2 Deterministic Networking Specification**

### **2.3.3 Traffic Shaping**

# Chapter 3

## Approach

To discuss the solution implemented for this thesis requires a review of, firstly, the requirements contained in the problem statement, and, secondly, a description of both the decisions made as well as the process behind making those decisions. To that extent, this chapter will discuss, the components needed in order to address the problem statement. It will try to discuss not only what decisions were made, but also, crucially, what decisions were *not* made, and why. The first section will begin with the fundamental nature of the solution's architecture, thereafter a review of the design requirements which arise from the Deterministic Networking specification will follow, finally the precise nature of the solution's internal workings will be presented.

### 3.1 Skeleton Structure

At the outset of this thesis it was decided to use a Control User Plane split. The packet processing and forwarding is performed in the user plane, and the control plane makes the high level decisions about which packets to send where, and in this case, on which outgoing interface to send them. The data plane is kept simple and only performs the time critical

packet processing, and blindly follows the instructions of the control plane. This type of architecture is very common in modern software-based networks, for example OpenFlow, the Evolved Packet Core (from LTE), and the 5G core all implement this kind of Control-Data split.

However, this decision immediately locks one into certain positions. For example, this requires a method of communication between the data plane and the control plane, it also means the state machine that represents the WAN Connector becomes more complex, and complexity always brings with it a greater danger for hidden mistakes and fallacies. The benefit is that the user plane function does not need to perform any of its own decisions, thus simplifying its internal architecture, and these components can be deployed at different locations, and scaled up or out with greater ease.

## 3.2 Architectural Components Required for Deterministic Multipath Backhaul in a 5G Campus Environment

Deterministic networks have been discussed in the previous section. In order to guarantee determinism the IETF DetNet working group has proposed an architecture for determinism over IP networks [**detnet-arch**]. Their specification identifies four key mechanisms for guaranteeing the determinism of a flow: 1) elimination of contention loss, 2) jitter reduction, 3) service protection, 4) explicit routes. Since the 5G campus environment may need to use the infrastructure of other operators this rules out the usage of explicit routes. A key difference between deterministic networks and 5G campus backhaul networks is that the operator may not control all the links between nodes, and that there are only two nodes in the network that are definitely under the administrator's control- the WAN Connectors at the core and the edge. In such a scenario the problem is less like a network and more like a client - server application with multiple paths between the client and the server.

### 3.2.1 Elimination of Contention Loss

Elimination of contention loss can be achieved by using a traffic shaper and/or rate limiter, and the ingress to any DetNet domain **must** apply such a function. For the WAN Connector therefore a traffic shaper must also be applied to the traffic on ingress, from the RAN, before it is sent across the backhaul links. Since the implementation of a traffic shaper is beyond the scope of this thesis, a choice must be made based on the existing solutions. The traffic control subsystem in the linux kernel (TC) provides several implementations of different algorithms for both rate limiting and traffic shaping. Some of these algorithms have already been discussed in the previous section. For the purposes of this solution several algorithms were considered: Hierarchical Token Bucket (HTB), Hierarchical Fair Service Curve (HFSC), Time Aware Priority Queueing (taprio), Common Applications Kept Enhanced (CAKE), and Fair Queueing with Controlled Delay (FQ-CoDel).

For HTB and HFSC, due to the implementation of these algorithms in the linux tc subsystem, integrating them within the context of this solution would require filters for each flow. This is not necessarily a hindrance yet, however it would introduce additional overhead upon the creation and/or addition of each new flow. Perhaps more importantly, the implementation of the HTB algorithm provides no means for fair queueing. This means that while it is not possible for greedy and/or malicious flows to use up the bandwidth of other flows because HTB guarantees bandwidth, the greedy flows can cause other flows to experience higher latency because it does not guarantee delay. The HFSC algorithm improves on this by offering both bandwidth and delay guarantees. However this makes it far more difficult to configure, and would require precise calculation of the service curve for each new flow.

Another feasible option available in linux TC would be the time aware priority queueing (taprio). This queueing discipline provides scheduled gates for specific traffic classes. It could be used to schedule different flows, and/or different priorities, and thus provide a form of prioritization, and, to an extent, traffic shaping. Flows can be assigned different windows (or

assigned windows based on their priority) and congestion due to a greedy flow is avoided. However one issue is that taprio fails to provide fair queuing, and thus would need to be configured and possibly adapted as new flows are added and removed, depending on their priority. Greedy flows of the same priority would have to fight each other to use the bandwidth available in their assigned window.

Lastly there are FQ\_CoDel and CAKE. Both provide fair queuing as well as active queue management, which means flows are less likely to experience loss due to congestion, particularly congestion caused by one particularly greedy flow. Since the CAKE algorithm was originally designed for routers in WLAN networks, and these are loosely analogous to a 5G campus network (consider the User Equipment as the Stations, and the Base Station as the Access Point), it makes it an attractive choice. CAKE also allows the use of priority bins, which fits well with the nature of 5G traffic, where different flows have different priorities. Finally, on a practical level, the implementation of CAKE in the linux kernel uses the kernel's "skb\_flow\_dissector" which exposes a hook point for eBPF [eBPF] programs. Specifically within a 5G context, this could allow one to attach an eBPF program to allow CAKE to peek inside of GTP tunnels, and differentiate the flows within them. For all the other algorithms mentioned so far, all of the GTP packets would appear to belong to the same flow, and thus none of these algorithms would work at all if the traffic is arriving in a GTP tunnel. This is ultimately the strongest case that can be made for any of the options listed so far. TC does have one other algorithm which hooks into the "skb\_flow\_dissector", namely the Choose and Keep Scheduler (CHOKe), which does perform queue management, however fails to provide fair queuing.

### 3.2.2 Jitter Reduction and Latency Guarantees

In the Deterministic Networking Architecture specification it is recommended to adopt time synchronization as well as sending "Time of Execution" fields in the application packets, in



order to achieve jitter reduction. For time synchronization, the PTP protocol may serve well, especially since there is an existing implementation- the linuxptp project - which extends it to work over IP networks. As for the time of execution fields, it may be possible to place these in a GTP header, or to combine them with timestamping.

Latency guarantees are not mentioned within the specification but they are important for the 5G campus setting. Specifically when backhaul options may include satellite links it becomes important to consider latencies. Additionally, in a multipath setting, latency can be a useful criteria for selecting between links, and guaranteeing latency becomes easier to do when it is possible to switch links should one of them start to experience greater latency than before.

### 3.2.3 Service Protection

In order to protect against equipment failure, it may be recommendable to perform packet duplication and/or encoding. The DetNet specification speaks of both duplication of flows, and network coding [**network-coding**]. This can be a clever solution, however Network Coding requires control over and co-ordination between all the nodes in the network (like other parts of the Deterministic Networking specifications). Though to guard against equipment failure and/or packet loss, duplication does provide an option. So too, does forward error correction (FEC). Many FEC schemes work on a continuous stream of bytes, providing correction for bit errors, as opposed to correcting the loss of entire packets. There are schemes which address this, fountain codes - such as Raptor [**raptor**], but unfortunately their ecosystem is not as developed and because of the complexity involved in implementing them efficiently there are relatively few freely available projects or libraries which can be used to encode packets using such a scheme. For reference, <sup>1</sup> maintains a list of C/C++ FEC libraries and none of them support the Raptor family of fountain codes (all other fountain codes incur high overhead). Even if it were more feasible to integrate FEC into the backhaul solution it bears questioning how much benefit it brings. FEC is very powerful in situations

---

<sup>1</sup><https://aff3ct.github.io/fec-libraries.html>

with consistently lossy links, however internet links tend to experience loss in bursts, as opposed to a consistent rate. To this extent it may make more sense to duplicate those flows which require a high degree of reliability. Flow duplication will mean a packet ordering and elimination function will be required, to eliminate those packets which arrive twice, and to re-order other ones.

If one does choose to duplicate packets in order to achieve service protection, that means an elimination function is also required, as well as a packet ordering function. This is because multiple packets may arrive on one link before they arrive on the backup path. If a packet got lost on the faster path, it may still show up on the slower one and thus the packet forwarding needs to pause until the missing packet arrives. The DetNet specifications provide both a basic and an advanced Packet Ordering Function (POF) algorithm. This algorithm can be very easily adjusted to also perform the elimination of duplicate packets. The DetNet specification also clarifies how long one should wait to see if the missing packet arrives (or doesn't) - this value "cannot be smaller than the delay difference of the paths used by the flow" and is called the *POFMaxDelay*.

```

/* initialization */
1 foreach flow  $f$  do  $POFLastSent_f \leftarrow 0$ ;
/* Start POF logic */
2 while true do
3   receive packet  $p$ , with sequence number  $seq\_num$  for flow  $f$ ;
4   if  $seq\_num \leq POFLastSent_f + 1$  then
5     /* eliminate duplicates */
6     if  $seq\_num == POFLastSent_f + 1$  then
7       forward  $p$ ;
8        $POFLastSent_f = seq\_num$ ;
9     end
10  else
11    /* in a separate thread */
12    buffer  $p$  until  $seq\_num = POFLastSent_f + 1$  or  $POFMaxDelay_f$  elapses;
13    forward  $p$ ;
14     $POFLastSent_f = seq\_num$ ;
15  end
16 end

```

**Algorithm 1:** Basic POF Algorithm Adjusted for De-Duplication

The algorithm shown in 1 is the Basic POF algorithm, with a single adjustment made on line 5 to make sure that an incoming packet's sequence number is not lower than the next one we are expecting, in order to properly eliminate duplicates. The algorithm's **if/else** logic in this representation could be unified in a cleaner way, but it was done this way specifically because it thus requires just one extra line (line 5) to be changed from the DetNet specifications POF algorithm, and that one change suffices to make it a Packet Ordering and Elimination Function.

### 3.2.4 Multipath Considerations

While the previous subsections have all considered requirements for determinism in an IP network, the issue of multipathing still remains. For a component which is backhauling over multiple links in an IP networks this means link and/or path selection is required. Traditionally, routers select links primarily based on their ability to route the packet to the destination, and secondarily based on various metrics, which can be defined by the administrator. For the WAN Connector's scenario routing considerations are not made, its purpose is to backhaul the traffic from the RAN / Edge to the Core, where the network's own internet gateway can perform the routing.

At this point it is crucial to differentiate, again, between multihoming and multipathing. Multihomed hosts are able to select one of many links for their outgoing traffic destined to the same source. In multipathing the same link may lead to multiple paths to *different* destinations. Multipathing over the internet requires complex data collection about all the routers in between the source and the destination, see [**multipath**] for an example of such an implementation. This is because multipathing for reliability over the internet needs to choose completely edge-disjoint paths to achieve maximum reliability.

The goal of this thesis is of course to use the link and/or path selection to provide determinism for the flows being backhauled. To this extent then, the link selection algorithm must attempt to forward flows over paths which can provide at worst the maximum allowed latency and jitter, as well as meeting the minimum reliability. Here it is worth noting that while duplication cannot really be used to guarantee latency, it can be used to improve reliability since a flow which is duplicated across two paths is far less likely to experience packet loss.

### **3.2.5 Summary of Requirements**

Looking back on the previous subsections, the following points (in no particular order of priority) were identified as mandatory for a deterministic backhaul solution: 1) traffic shaping, 2) path selection according to jitter and latency requirements, 3) service protection (i.e. via packet duplication), 4) packet ordering and de-duplication, and 5) time synchronization. The ways in which these can be addressed, or quite simply the way in which they were implemented, will be discussed in the following section

## **3.3 Overview of the WAN Connector's Features and Components**

At this point, one can take the skeleton structure proposed in the first section, and superimpose the other requirements which were determined in the second section. Combining these ideas yielded the final implementation, which will be now be presented.

### **3.3.1 Path Selection Algorithm**

Meeting the various requirements - jitter, latency, and delay - of a flow can be formulated as a multi-constrained QoS problem. Solving such a multi-constrained QoS problem via path selection is a binary optimization problem. The problem can be posed as: "select those paths on which to forward packets while making sure to satisfy the latency, jitter and reliability requirements of the given flow, and minimizing the overall weight of the paths used". The mathematical definition is as follows:

$$\text{Minimize } \sum_{i=1}^P w(x_i) \quad (3.1)$$

$$\text{Where, } d(i) * x_i \leq D \quad (3.2)$$

$$j(i) * x_i \leq J \quad (3.3)$$

$$1 - \prod_{i=1}^P (1 - r(i) * x_i) \geq R \quad (3.4)$$

$$\text{for } x_i \in \{0, 1\} \quad (3.5)$$

Here the variables  $D$ ,  $J$ , and  $R$  are the flow's delay, jitter, and reliability requirements, while the functions  $d(i)$ ,  $j(i)$ ,  $r(i)$ ,  $w(i)$  are the estimated delay, jitter, reliability, and weight of link  $i$ . The predicted values will usually just be the latest measurement, as recommended in [2], however there is room here to use more advanced metrics to predict the future link quality and thus perform preemptive path switching in future work. The total number of paths is  $P$ . The  $x_i$  variable indicates whether or not link  $i$  shall be used. If a solution is found, then the flow's packets will be forwarded on each link  $i$  where  $x_i = 1$ , and if no solution can be found which satisfies these conditions then the flow is rejected because its QoS cannot be guaranteed.

It is worth noting that solving such problems is NP-Hard [**tsp-np-hard**]. However this hardness arises primarily because of equation 3.4, the equation for reliability (also the only non-linear equation). Due to this equation one must consider every possible combination of paths on which to forward, and the complexity is  $O(2^n)$ . This means, even though akella et. al [1] have shown that multihomed approaches experience diminishing returns after more than 4 links, attempting to brute force the solution by limiting the number of outgoing interfaces to 4 still yields a very large problem space - in the worst case both WAN connectors could have 4 outgoing paths, leading to 16 possible paths and thus 65536 possible combinations to consider.

In order to further avoid the combinatorial explosion the problem needs to be parameterized even more. The first logical parameterization has already taken place by limiting the number of interfaces to 4. This can be expanded on by limiting duplication to only take place over disjoint interfaces. The reasoning behind this is that in a geographically distributed Campus 5G environment, and especially for environments featuring wireless backhaul (e.g. satellite), it is more likely that a link’s reliability is most affected by the over the air transmission on the first hop. By performing this parameterization the problem space shrinks considerably, to just  $2^4$  possible combinations of outgoing paths, but, because certain paths going out on the same interface are ignored, the optimality of the overall solution is gone. This is an acceptable trade off for quick computation, and reasonably strong guarantees on reliability. When choosing which path to include from a set of paths using the same outgoing link, only the path with the greatest reliability is taken into consideration so that the likelihood of rejecting a potentially viable flow is as low as possible.

### 3.3.2 Packet Ordering and De-Duplication Function

Since it is possible for flows to be duplicated across multiple paths, it becomes a necessity to have a packet ordering and elimination function. To be able to re-order and de-duplicate packets means that their sequence numbers need to be tracked. This thesis’ implementation will track the sequence number on a per flow basis, using the GTP ”sequence number” field. This thesis uses it’s variation of the DetNet POF specification’s Basic POF algorithm [[detnet-pof](#)], as presented in [1](#).

### 3.3.3 Internal Architecture

#### Control - Data Plane Communication

A protocol is desired which can provide reliable communication over multiple paths in order to communicate between the control plane, where the statistics are collected, flows are added or removed, and the multipath decisions are made, and the data plane. This is crucial, since in a geographically distributed campus 5G deployment it is unlikely that there will be a separate management or control network which uses different underlying network infrastructure than the data plane. Therefore link failure, as well as packet loss, on the link being used by the control plane must either be avoided or protected against. To this extent there are only two feasible options - either multipath TCP or SCTP. SCTP does support multihoming and intelligent failover, despite not being explicitly designed for multipathing, and is slightly easier to manage and configure than multipath TCP, which is why it was chosen for the data plane control plane interactions.

#### GTP Tunneling and Custom Header Extension

The tunneling protocol used between the two data plane instances of the WAN Connector will be GTP version 1 [**GTP-spec**]. The specification for GTP allows for the use of sequence numbers, as well as the use of extension headers. For the data plane communication there will always be an additional custom header sent which includes a timestamp taken by the sender. The timestamp is taken in microseconds. For the timestamp there are several options available within linux. The TAI clock in linux (representing the Atomic International Time) is used in this implementation because it does not have leap seconds, and so is a monotonic function, which is an important guarantee for time sensitive applications.

Normally, in 64 bit linux systems, a timestamp takes up 16 bytes and consists of 8 bytes for the time in seconds since the epoch, and another 8 bytes for the nanoseconds. To reduce the



footprint of the timestamp in the header, the seconds are represented with just 8 bytes, thus wrapping around the interval  $[0, 255]$ . The receiver needs to take this into account, but is programmed to do so, and since the path delays are not realistically expected to ever exceed 3 seconds this is more than enough. The microseconds have a maximum value of  $10^6$  and can be represented with just 24 bits. For Wide Area Networks and internet connections usually the latencies are on the order of milliseconds and as such microsecond precision is deemed to be sufficient.

## **Metric Collection**

To be able to intelligently switch flows between paths, and to be able to know when this is required, it is necessary to collect the relevant metrics about latency, jitter, bandwidth usage, as well as packet loss. Packet loss can be detected via the sequence numbers, while delay and jitter can be calculated using the timestamps passed along in the GTP headers. These metrics are periodically reported to the control plane so that it can make its decisions on up to date data. Keeping a healthy overview over the state of each path requires periodic probing on these paths. This is especially important for detecting when paths become viable again, since these paths will not have any traffic on them while they are considered down, or if they have experienced high latency and/or jitter recently. The probes are sent once per period of reporting so that they have a minimal impact on the bandwidth usage.

## **Flow Descriptions, Flows, and Hashing**

Incoming packets need to be quickly mapped to their respective flows. It is common in network environments to perform flow hashing, in order to quickly lookup which flow incoming packets belong to. This approach makes sense here too. However, since 5G flow descriptions can apply to various IP and port ranges as well as matching based on the transport layer protocol it is not possible to hash an incoming flow and obtain the same hash as the flow

description. Furthermore it is possible for a flow to match multiple different flow descriptions. In these cases usually the first matching flow description is taken. This implies some sort of ordered storage of flow descriptions will be required, as well as a method to quickly match packets to flows. The solution used here is to maintain a simple linked list of the known flow descriptions, and match packets to their flow descriptions, if the packet is unknown. For "known" packets, which have been matched to a flow description, these are hashed and stored in a table, alongside their flow descriptor and a list of paths on which the flow is supposed to be forwarded. This allows quick lookup for every subsequent packet, as well as making it simple to change which paths flows are meant to be forwarded on. One alternative to storing the flow descriptions in a linked list would be to store them in a tree, such as in [flow-lookup-trees-phd], sorted either by priority or ID or even by the hash value of the flow description. This method would provide a faster lookup of new, "unknown" packets, whose hashes don't yet match to a flow, but since new flows are not such a frequent event the overhead may not be worth the gain in performance.

It is important to make a good choice of a hash function. Since the hashes used are only for the purpose of lookups it makes no sense to pick a hashing algorithm designed for use in cryptography. These algorithms are designed so that it is very difficult to reverse engineer the original value from the hash, and this may often make the computation of the hash more computationally intensive than when computing hash algorithms designed for looking up table entries. Lastly, for hashing algorithms it is desirable that they provide a healthy distribution, to reduce collisions. Collision reduction can also be affected by choice of hash table size, in general it is usually recommended to use prime number. In this implementation, the MurmurHash algorithm was chosen due to its strong performance for lookup-based hashing. Specifically, the MurmurHash3 [6] version was chosen, and since it can generate 32 or 128 bit values - the exact sizes of IPv4 and IPv6 addresses - this simplifies the hashing implementation for IP flows. When a packet comes into the WAN connector it is hashed based on it's destination address, source address, destination port, source port, and finally the transport layer protocol number. If the packet was tunneled with a GTP header, the

header is removed and the hash is performed on the tunneled packet, not on the GTP packet. If the hash does not match, the packets is compared to the list of existing flow descriptions. If the flow belongs to one of these descriptions then that decision is stored in the hash table.

### 3.3.4 Overview

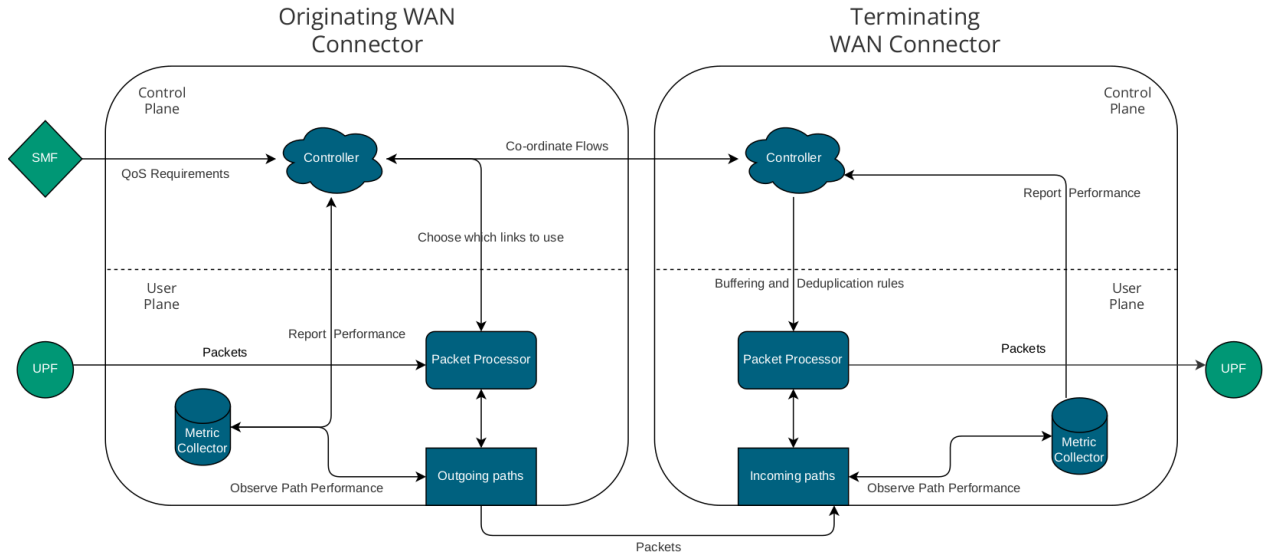


Figure 3.1: Internal Architecture of the WAN Connector

The internal architecture of the WAN connector is shown in Figure 3.1. The connector is split into a control plane, which intelligently chooses which interfaces to forward on, and a data plane, which performs the forwarding, as well as any necessary buffering, re-ordering and/or de-duplication. The WAN connector also performs different tasks depending on whether it is the origin or termination of a flow. For example, the terminating node may be receiving duplicate packets on the other paths, and it must know to drop these. Additionally, the sender, can only store statistics about how many bytes it has sent, and on what paths. The receiver is able to use the information contained in the packets it has received to construct a complete picture about the nature of the path - its jitter, latency and packet loss. The link selector in the control plane can then use this information to make its decisions.

One part which is missing from this diagram is the traffic shaper. That is because it is not connected to the control and the data plane, rather it is meant to be installed alongside them. Furthermore, it is important that the traffic shaper resides in the part of the network before the packets enter the WAN connector. If the deployment features a UPF in the edge network, it may make sense to place the traffic shaper before this as well. For the core, the traffic shaper will similarly need to be placed at the ingress point at which packets reach the WAN connector.

# Chapter 4

## Evaluation

This chapter reviews how the approach from the previous chapter was tested, and it discusses the observed results. The goal here is to explain the setup of the testbed, as well as the motivation behind the types of tests which were performed. Finally a discussion of the results and their implications will take place.

### 4.1 Approach to Evaluating Performance

In order to evaluate the success of the proposed approach, multipath scenarios will be emulated and investigated. Each setup will consist of two WAN connectors, and some number of emulated links going between them. Due to the expectation that satellite backhaul will become more and more relevant for geographically distributed campus networks, the emulation will feature a Low Earth Orbit (LEO) satellite link as well as a Geostationsary Orbit (GEO) link, and a terrestrial connection, which will be based on the characteristics of DSL or perhaps earlier DOCSIS links, but specifically **not** optical fibre. This thesis is addressed towards campus network deployments that do not have access to fibre yet, and cannot expect to receive it in the near future.




Figure 4.1: Testbed Setup

In all of these scenarios the same traffic flows will be replayed. This traffic will contain different types of flows, with different QoS requirements. Before a new flow is started, the flow's requirements are sent to the WAN connector and it is either accepted or rejected. During the traffic replay, the delay, jitter, and reliability will be measured.

All of these tests can be performed on the same testbed which will be set up as is shown in Figure 4.1. The testbed architecture features a traffic generator, two WAN connectors, with a link emulator between them, and a measurement module to analyze performance. In practice the traffic generator and the test bench will be co-located on the same machine. The link emulator, and the WAN connectors are separate hosts, all interconnected over ethernet. The link emulation is done using the linux kernel's traffic control subsystem (TC). TC offers a network emulator (netem) queuing discipline (qdisc), which is able to emulate various link characteristics including delay, jitter, packet loss and packet re-ordering. Furthermore by combining the netem qdisc with a rate limiter, such as the Heirarchical Token Bucket (HTB) qdisc, a link can have its bandwidth limited. This allows one to emulate links with different latency, reliability and jitter, and with different maximum bandwidths.

In order to make the emulation more realistic, it will periodically be adjusted, this allows one to degrade a link over time, for example by increasing the latency or packet loss in frequent intervals, up to a large value. It also more closely mimics the real-life behavior of WAN connections, which do experience changes in their characteristics over time.

For the purposes of evaluating the WAN connector, two initial series of experiments will be performed to isolate and investigate its link switching capabilities. Firstly, purely latency based link selection will be investigated- a flow will be defined with specific latency requirements and then the emulated links will have their latency repeatedly changed so that one or more of them begin to violate the maximum latency. The WAN connector should then switch the flow to a different path. The same will be done once with packet loss. First, a link will be given a packet loss that is higher than the acceptable level for the given flow. Then multiple links will be adjusted so that their level of packet loss is too high to support a flow on its own, later complete link failure will be simulated. This simulation will be done by using the linux packet filtering tool iptables to block all traffic with the source IP of one of the links. These two experiments will both be conducted with other variables held to a control value. So first only the latency will be varied, then only the packet loss. In the experiments with packet loss the WAN connector should at first switch the flow to a different connection, and then later being to duplicate the flow. During the duplication experiment, the application behind the flow should not receive any duplicate packets.

Lastly after the two control experiments have confirmed (or not confirmed) the expected behaviour, a proper test will take place where the links' characteristics are varied independently of each other and once again a select few critical flows will be observed, to make sure the WAN connector is able to uphold their QoS requirements.

## **4.2 Latency Based Path Switching**

## **4.3 Reliability Based Path Switching**

## **4.4 Extended Experiments**





# Chapter 5

## Conclusion and Outlook

Looking back, this thesis has addresses the usage of multiple backhaul paths in a campus 5G setting to achieve determinism. A review of the topic and relevant literature was done in the background chapter, then an approach was developed based on that information and based on the requirements of the problem. Finally this approach was evaluated in a testbed with a link emulator to see how it performed.

TODO

### 5.1 Improvements

There are several possible improvements that immediately spring to mind. Firstly, the path estimation could be adjusted to not just report previous statistics but also attempt to infer what the path might look like in the near future, for example as a path begins to experience increased latency a predictive algorithm/approach might be able to pre-emptively move flows off of that path, before their latency requirements are violated. This approach could be based on machine learning or AI, or it could use analytical methods.

Additionally, specifically for the Lower Earth Orbit (LEO) satellite case, the path selection could potentially be adjusted to account for the periodic increases in latency as the current satellite leaves the range of the ground station, and the next one comes into range. For example during this phase it might make sense to temporarily forward packets on a different path.

Another potential improvement would be to pass a "Time of Execution" field in the GTP header of packets of jitter-sensitive applications. This allows the receiving WAN connector to hold packets if they have arrived too early, and, conversely, the packet ordering function may use this field to determine that it is more important to forward the current packet now, than to wait for a missing packet.

As the authors in [adaptive] did, this thesis' approach could benefit from adaptive windows of reporting. This way there is not additional overhead with overly frequent statistical updates, as well as avoiding the reverse situation, where the reporting is too infrequent for a rapidly changing path.

The last possible improvement is the addition of Forward Error Encoding (FEC). While this would be difficult to integrate into the equation to select paths, FEC could be used to increase the resilience of consistently lossy links, which is a big benefit for links which commonly exhibit this characteristic, such as wireless links.

## 5.2 Implications and Further Areas of Research

Since the results were only verified in a testbed setup it would make sense to now test the WAN Connector in a real campus 5G deployment where there actually are multiple outgoing paths.

Research should also be conducted on evaluating the performance of specific applications performance. For example the quality of VoIP calls do not only depend on latency and

jitter [**voip-measurement**], but rather how they interact together, and they have their own suites of evaluation criteria. Another specific application to consider has to be interactive video. Mission critical as well as control systems could also be pulled into the test suites for evaluation.

This approach has been a deterministic one, that is the function used to select paths on which to forward is deterministic. It would be interesting to see what benefit AI and machine learning approaches may bring to this problem since they can act more dynamically, and perhaps learn the characteristics of a given link over time. Perhaps they can discover what the link exhibits right before total failure and thus perform pre-emptive path switching.



# Bibliography

- [1] Aditya Akella et al. “A measurement-based analysis of multihoming”. In: *Proceedings of the 2003 conference on Applications, technologies, architectures, and protocols for computer communications*. 2003, pp. 353–364.
- [2] Aditya Akella et al. “On the performance benefits of multihoming route control”. In: *IEEE/ACM Transactions on Networking* 16.1 (2008), pp. 91–104.
- [3] Hind Alwan and Anjali Agarwal. “Multi-objective reliable multipath routing for wireless sensor networks”. In: *2010 IEEE Globecom Workshops*. IEEE. 2010, pp. 1227–1231.
- [4] Jennifer Andreoli-Fang and John T Chapman. “Mobile-aware scheduling for low latency backhaul over DOCSIS”. In: *2017 IEEE 28th Annual International Symposium on Personal, Indoor, and Mobile Radio Communications (PIMRC)*. IEEE. 2017, pp. 1–6.
- [5] Maria Apostolaki, Ankit Singla, and Laurent Vanbever. “Performance-driven internet path selection”. In: *Proceedings of the ACM SIGCOMM Symposium on SDN Research (SOSR)*. 2021, pp. 41–53.
- [6] Austin Appleby. “MurmurHash3, 2012”. In: *URL: <https://github.com/aappleby/smhasher/blob/master/src/murmur3.cpp>* (2012).
- [7] Xavier Artiga et al. “Terrestrial-satellite integration in dynamic 5G backhaul networks”. In: *2016 8th advanced satellite multimedia systems conference and the 14th signal processing for space communications workshop (ASMS/SPSC)*. IEEE. 2016, pp. 1–6.

- [8] Anat Bremner-Barr et al. “Predicting and bypassing end-to-end Internet service degradations”. In: *Proceedings of the 2nd ACM SIGCOMM Workshop on Internet Measurement*. 2002, pp. 307–320.
- [9] Nelson Capela and Susana Sargento. “Multihoming and network coding: A new approach to optimize the network performance”. In: *Computer Networks* 75 (2014), pp. 18–36.
- [10] Kameswari Chebrolu and Ramesh R Rao. “Bandwidth aggregation for real-time applications in heterogeneous wireless networks”. In: *IEEE Transactions on Mobile Computing* 5.4 (2006), pp. 388–403.
- [11] Fuqiao Chen and Chenyang Yin. “A Collaborative Traffic Scheduling Mechanism for Multi-homed Networks”. In: *2020 IEEE 5th Information Technology and Mechatronics Engineering Conference (ITOEC)*. IEEE. 2020, pp. 177–181.
- [12] Joerg Deutschmann, Kai-Steffen Hielscher, and Reinhard German. “Broadband internet access via satellite: Performance measurements with different operators and applications”. In: *Broadband Coverage in Germany; 16th ITG-Symposium*. VDE. 2022, pp. 1–7.
- [13] Doğanalp Ergenç and Mathias Fischer. “On the reliability of ieee 802.11 cb frs”. In: *IEEE INFOCOM 2021-IEEE Conference on Computer Communications*. IEEE. 2021, pp. 1–10.
- [14] Norman Finn and Pascal Thubert. *Deterministic networking problem statement*. Tech. rep. 2019.
- [15] Norman Finn et al. “Deterministic networking architecture”. In: *RFC 8655*. IETF, 2019.
- [16] Igor Ganichev et al. “YAMR: Yet another multipath routing protocol”. In: *ACM SIGCOMM Computer Communication Review* 40.5 (2010), pp. 13–19.
- [17] David K Goldenberg et al. “Optimizing cost and performance for multihoming”. In: *ACM SIGCOMM Computer Communication Review* 34.4 (2004), pp. 79–92.

- [18] Fanglu Guo et al. “Experiences in building a multihoming load balancing system”. In: *IEEE INFOCOM 2004*. Vol. 2. IEEE. 2004, pp. 1241–1251.
- [19] Ahsan Habib and John Chuang. “Improving application QoS with residential multihoming”. In: *Computer Networks* 51.12 (2007), pp. 3323–3337.
- [20] Toke Høiland-Jørgensen, Dave Täht, and Jonathan Morton. “Piece of CAKE: a comprehensive queue management solution for home gateways”. In: *2018 IEEE International Symposium on Local and Metropolitan Area Networks (LANMAN)*. IEEE. 2018, pp. 37–42.
- [21] Chung-Ming Huang and Ching-Hsien Tsai. “WiMP-SCTP: Multi-path transmission using stream control transmission protocol (SCTP) in wireless networks”. In: *21st International Conference on Advanced Information Networking and Applications Workshops (AINAW’07)*. Vol. 1. IEEE. 2007, pp. 209–214.
- [22] Xiaoxia Huang and Yuguang Fang. “Multiconstrained QoS multipath routing in wireless sensor networks”. In: *Wireless Networks* 14.4 (2008), pp. 465–478.
- [23] Mona Jaber et al. “5G backhaul challenges and emerging research directions: A survey”. In: *IEEE access* 4 (2016), pp. 1743–1766.
- [24] Ralf Kundel et al. “User Plane Hardware Acceleration in Access Networks: Experiences in Offloading Network Functions in Real 5G Deployments.” In: *HICSS*. 2022, pp. 1–10.
- [25] Stanislav Lange et al. “Performance benchmarking of a software-based LTE SGW”. In: *2015 11th International Conference on Network and Service Management (CNSM)*. IEEE. 2015, pp. 378–383.
- [26] Ming Li et al. “Multipath transmission for the internet: A survey”. In: *IEEE Communications Surveys & Tutorials* 18.4 (2016), pp. 2887–2925.
- [27] Thinh Nguyen and Avidah Zakhor. “Path diversity with forward error correction (pdf) system for packet switched networks”. In: *IEEE INFOCOM 2003. Twenty-second Annual Joint Conference of the IEEE Computer and Communications Societies (IEEE Cat. No. 03CH37428)*. Vol. 1. IEEE. 2003, pp. 663–672.

- [28] Jonathan Prados-Garzon and Tarik Taleb. “Asynchronous time-sensitive networking for 5G backhauling”. In: *IEEE Network* 35.2 (2021), pp. 144–151.
- [29] Pablo Rodriguez et al. “Mar: A commuter router infrastructure for the mobile internet”. In: *Proceedings of the 2nd international conference on Mobile systems, applications, and services*. 2004, pp. 217–230.
- [30] Reza Sheyibani et al. “A reliable and qos aware multi-path routing algorithm in wsns”. In: *2012 Third International Conference on Emerging Intelligent Data and Web Technologies*. IEEE. 2012, pp. 125–132.
- [31] Ion Stoica, Hui Zhang, and TS Eugene Ng. “A hierarchical fair service curve algorithm for link-sharing, real-time and priority services”. In: *ACM SIGCOMM Computer Communication Review* 27.4 (1997), pp. 249–262.
- [32] Shu Tao and Roch Guérin. “Application-specific path switching: A case study for streaming video”. In: *Proceedings of the 12th annual ACM international conference on Multimedia*. 2004, pp. 136–143.
- [33] Shu Tao et al. “Exploring the performance benefits of end-to-end path switching”. In: *Proceedings of the joint international conference on Measurement and modeling of computer systems*. 2004, pp. 418–419.
- [34] Shu Tao et al. “Improving VoIP quality through path switching”. In: *Proceedings IEEE 24th Annual Joint Conference of the IEEE Computer and Communications Societies*. Vol. 4. IEEE. 2005, pp. 2268–2278.
- [35] Mohammed Tarique et al. “Survey of multipath routing protocols for mobile ad hoc networks”. In: *Journal of network and computer applications* 32.6 (2009), pp. 1125–1143.
- [36] Jack Tsai and Tim Moors. “A review of multipath routing protocols: From wireless ad hoc to mesh networks”. In: *ACoRN early career researcher workshop on wireless multihop networking*. Vol. 30. Citeseer. 2006.



- [37] Simon Tschöke et al. “Time-sensitive networking over metropolitan area networks for remote industrial control”. In: *2021 IEEE/ACM 25th International Symposium on Distributed Simulation and Real Time Applications (DS-RT)*. IEEE. 2021, pp. 1–4.
- [38] Zhiqiang Xiong et al. “A lightweight FEC algorithm for fault tolerant routing in wireless sensor networks”. In: *2006 International Conference on Wireless Communications, Networking and Mobile Computing*. IEEE. 2006, pp. 1–4.
- [39] Pouria Zand et al. “Wireless industrial monitoring and control networks: The journey so far and the road ahead”. In: *Journal of sensor and actuator networks* 1.2 (2012), pp. 123–152.
- [40] Gongzheng Zhang et al. “Fundamentals of heterogeneous backhaul design—Analysis and optimization”. In: *IEEE Transactions on Communications* 64.2 (2016), pp. 876–889.